



Data Science Initiative
BROWN



Final Project Presentation

— OYO rental price prediction in China

Brown University, Data Science Initiative, 22fall

GitHub: https://github.com/AstrosiosaurQ7/data1030_final_project.git

Xiaoyan Liu

06/12/2022



Content

1

Recap

2

Cross validation

3

Results

4

Improve



Data Resource

OYO APP , Chinese OTA platform :
XieCheng / MeiTuan/FeiZhu


Kaggle [Oyo Rental Price Prediction in China | Kaggle](#)

- This report hopes to predict the rental price of OYO hotels according to the property type 、 hotel location and so on.


Target variable : rental price / \$

Regression / Right-skewed

Feature matrix shape: (5834,25)




➤ Enlarge the coordinate axis by log function



Most hotels are concentrated in the surrounding cities of Shenzhen and Hong Kong, which are near to the seaport and have developed tourism and economic industries.

```
1 df[['longitude(East)', 'latitude(North)']]
```

	longitude(East)	latitude(North)
0	114.059600	22.542900
1	114.043225	22.539490
2	114.079426	22.508573
3	114.079035	22.508697
4	114.055590	22.509502
...
5829	114.194588	22.619618
5830	114.201327	22.606116



Step 1. drop worthless features : amenities, has_availability

Step 2. preprocessor encoding

- One-Hot encoder : 'bed_type', 'host_is_superhost', 'instant_bookable', 'room_type'
- Ordinal encoder : 'cancellation_policy', 'property_type'
- MinMax Scaler : 'accommodates', 'availability_30' , 'calculated_host_listings_count' , 'guests_included' , 'number_of_reviews', 'bathrooms', 'bedrooms', 'beds' , 'host_listings_count' , 'review_scores_checkin', 'review_scores_communication' , 'review_scores_location' , 'review_scores_rating', 'review_scores_value'
- Standard Scaler : 'maximum_nights'



def MLpipe_KFold_RMSE

- Split 20% data to be test set
- Used **Kfold**
(**n_splits=4, shuffle=True**)
- GridSearchCV function
- evaluation metric= **RMSE**
- Output best parameters and best score
- Save the scores and **y_test** for the baseline model

```
def MLpipe_KFold_RMSE(X, y, preprocessor, reg, param_grid):  
    test_scores = np.zeros(nr_states)  
    for i in range(nr_states):  
  
        X_other, X_test, y_other, y_test = train_test_split(X, y, test_size = 0.2, random_state=42*i)  
        kf = KFold(n_splits=4, shuffle=True, random_state=42*i)  
        pipe = make_pipeline(preprocessor, reg)  
        grid = GridSearchCV(pipe, param_grid=param_grid, scoring = 'neg_mean_squared_error',  
                             cv=kf, return_train_score = True, n_jobs=-1, verbose=True)  
        grid.fit(X_other, y_other)  
        results = pd.DataFrame(grid.cv_results_)  
  
        print('best model parameters:', grid.best_params_)  
        print('validation score:', grid.best_score_)  
  
        # save the model  
        final_models.append(grid)  
  
        # calculate and save the test score  
        y_test_pred = final_models[-1].predict(X_test)  
        test_scores[i] = np.sqrt(mean_squared_error(y_test, y_test_pred))  
        print('RMSE test score:', test_scores[i])  
  
    return test_scores, y_test
```



L1 regularized linear regression(Lasso)

- ✓ alpha : 0.25, 2.5, 25
- ✓ random_state=42

RandomForestRegressor

- ✓ max_features : [3, 5, 7, 9]
- ✓ max_depth : [3, 5, 7, 9]
- ✓ random_state=42

SVR (Support Vector)

- ✓ gamma: [0.1, 10, 100]
- ✓ C : [0.1, 1, 10]
- ✓ kernel='rbf'

KNeighborsRegressor(

- ✓ n_neighbors': [5, 25, 50]
- ✓ weights='uniform'

**L1 :**

- test score: [332.50481953, 321.1485038 , 249.33466631, 369.32089209, 298.27287924, 349.18486624, 364.13289666, 258.50704218, 229.58742343, 337.10134802]
- test scores mean : 310.9095337515326
- test scores standard deviation : 47.198

SVR :


- test score: [457.0634567, 439.1388, 352.7975, 502.4557654, 409.218457, 442.50976546yy, 495.1113467, 357.6800, 299.077, 460.71044]
- test scores mean : 421.576384
- test scores standard deviation : 62.8157611

RandomForestRegressor :

- test score: [307.10369093, 305.09515619, 215.86303583, 344.76157886, 292.349, 336.76020741, 360.37117433, 239.6237, 218.4528, 299.79961715]
- test scores mean : 292.01818453
- test scores standard deviation : 48.858289905

KNeighborsRegressor :

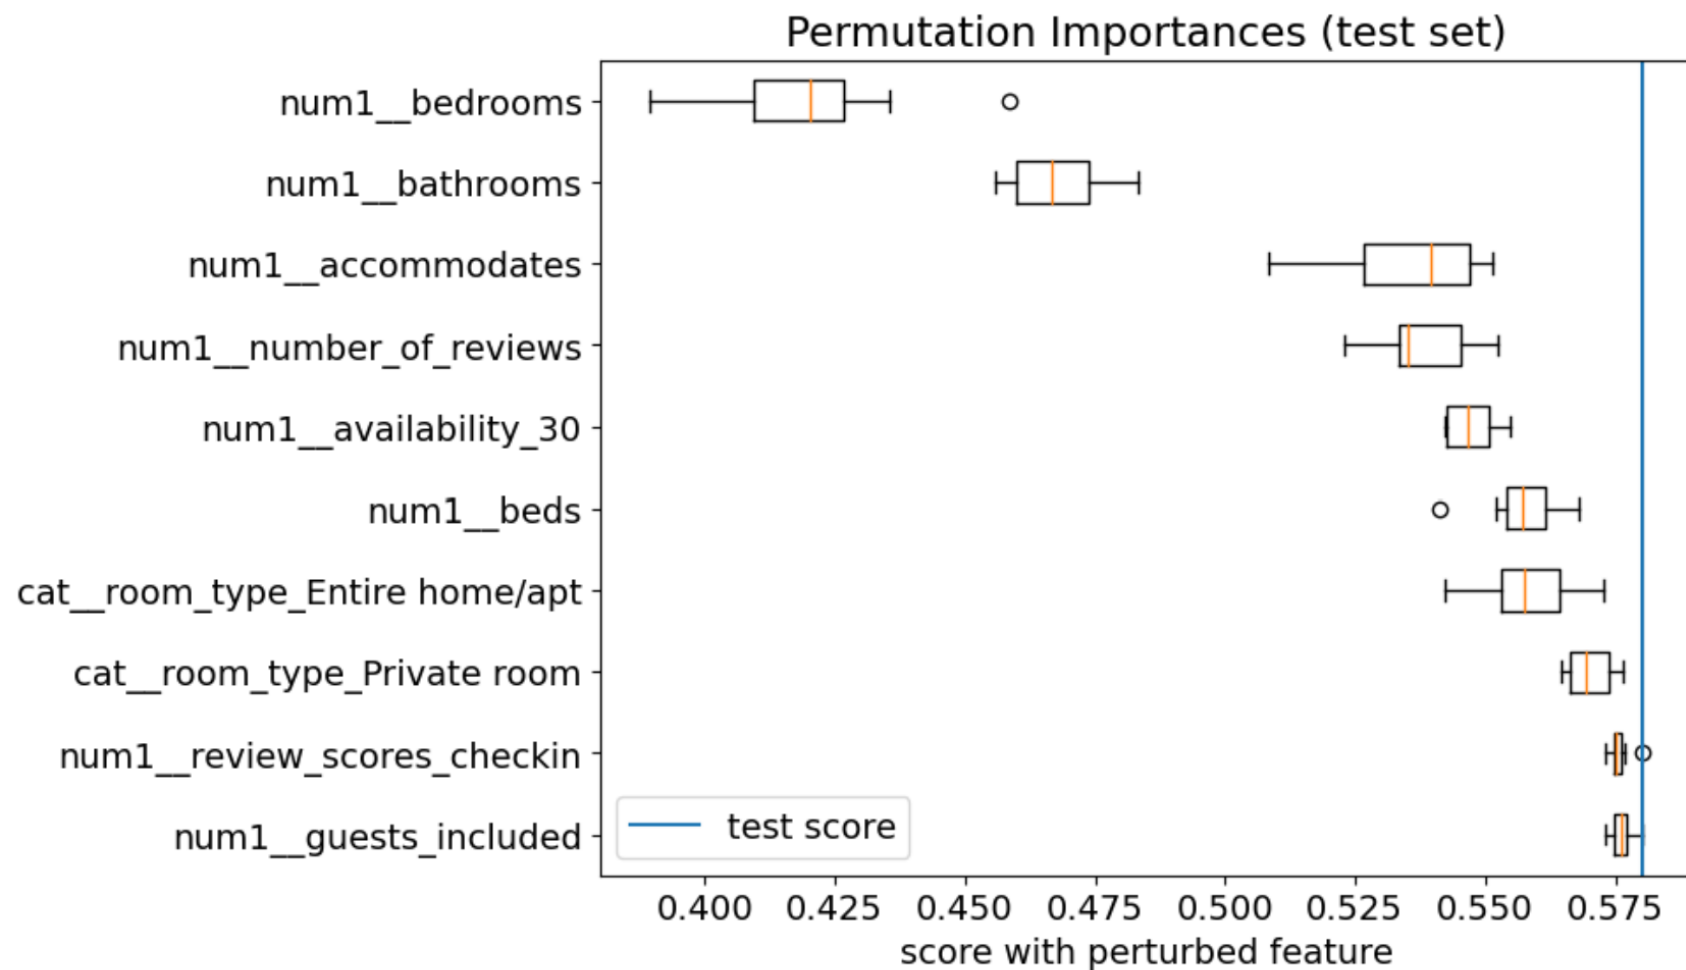
- test score: [318.20838074, 320.60138962, 265.89838554, 371.5695846 , 289.9438, 335.33792052, 384.54992622, 252.862433 , 222.87348307, 338.97971562]
- test scores mean : 310.082503
- test scores standard deviation : 49.206679





- baseline RMSE: 415.7852569807543
($y_{\text{mean}} = y_{\text{pred}}$)
- test scores mean : 333.64
- test scores standard deviation : 73.346
- best model parameters:
{'randomforestregressor_max_dept
h': 9,
'randomforestregressor_max_featur
es': 5}



- best model parameters: `{'max_depth': 9, 'max_features': 5}`
- train set shape : (4667, 30)
- test set shape : (583, 30)
- running time 0.3580458164215088
- test scores 202.72660072854498

**Shuffling****TOP 10**

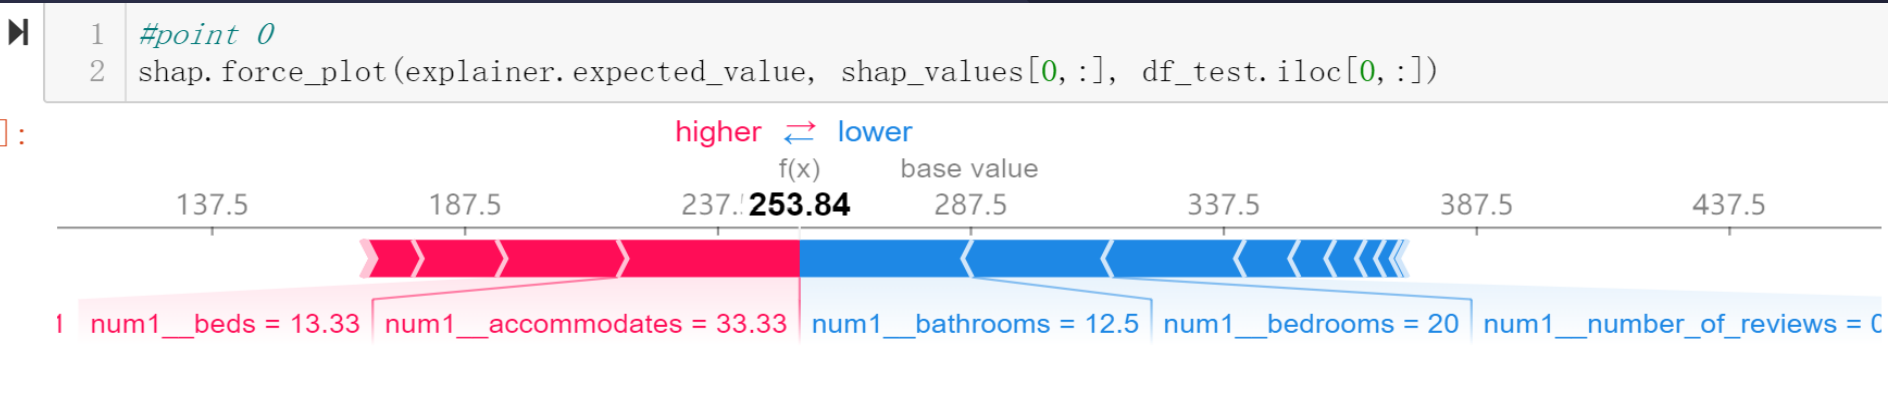
**SHAP Global****TOP 10**



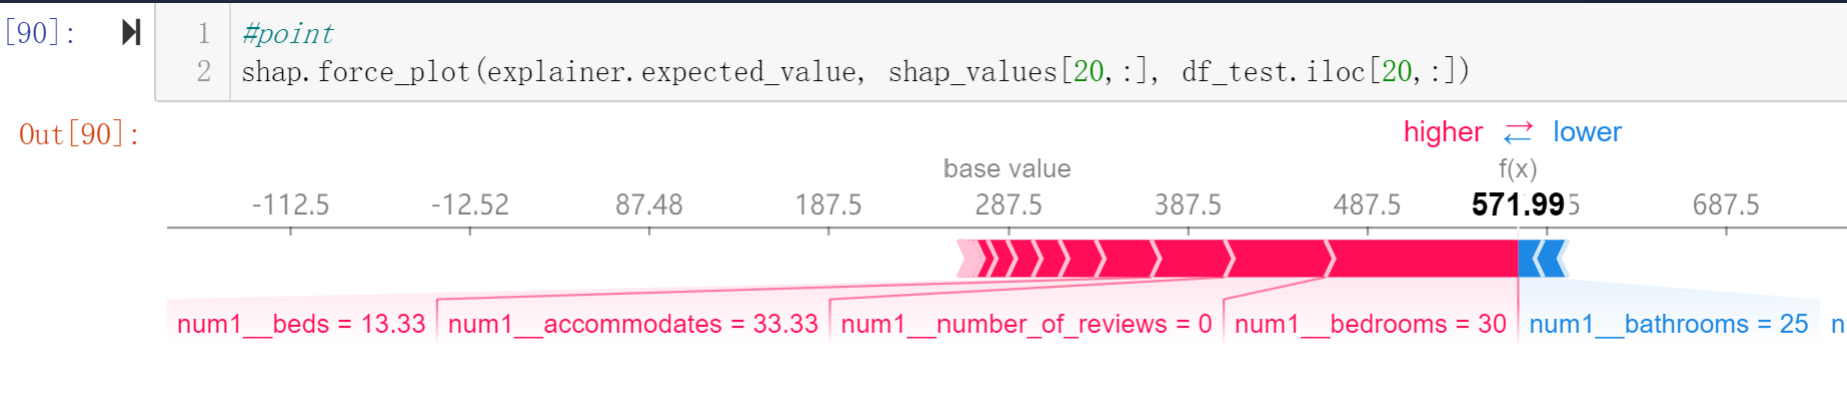
Randomforest feature_importanc es function

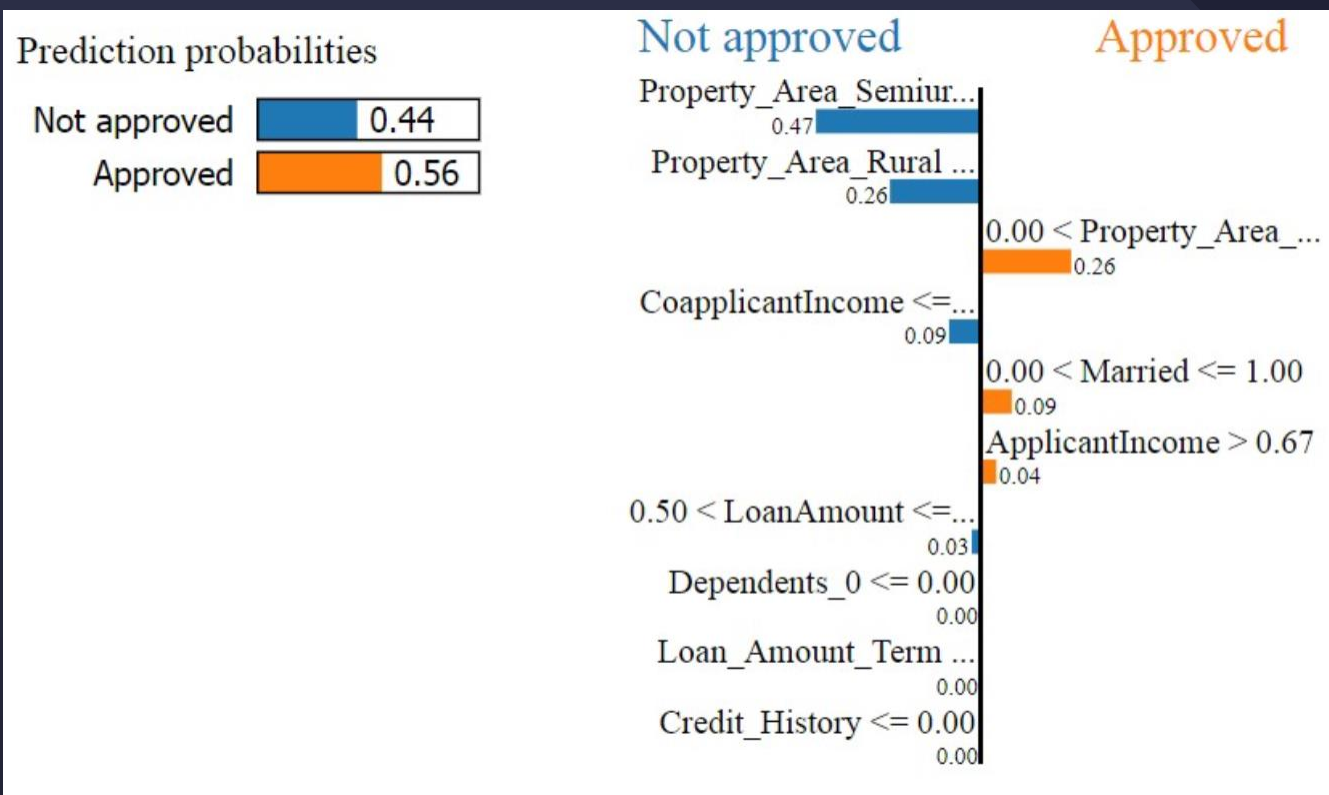
- Those results are very similar to the previous , but the sort is different.

Point 0



Point 20





- SHAP, LIME, and Anchors, provide local, model-agnostic interpretability methods.

Reference:

<https://towardsdatascience.com/three-interpretability-methods-to-consider-when-developing-your-machine-learning-model-5bf368b47fac>

<https://christophm.github.io/interpretable-ml-book/interpretability.html>

- XGBoost



Data Science Initiative
BROWN



THANK YOU FOR WATCHING

